

# Lessons Learned From Refactoring Multi Terabyte Databases

for the past 10 years



# Sponsors & Organizers



# What I Want to Talk About

**Refactoring** is a disciplined technique for restructuring an existing body of code, altering its internal structure without changing its external behavior.

Martin Fowler





Image Source: [https://www.regulation.org.uk/ob-failure\\_to\\_learn.html](https://www.regulation.org.uk/ob-failure_to_learn.html)

What was the biggest database  
you've ever worked with?

# About Me - Sergey Olontsev

Data Engineering Team Lead @ IPONWEB (a part of Criteo), ex-Kaspersky (Staff Engineer), ex-MVP.

18 years of experience with databases and high load systems (started with SQL Server 2000)

Worked with OLTP databases up to 50 TB, relational DWH up to 200 TB. Each, not in total! :)

100+ billions of rows in tables (OLTP).

<https://sergeyolontsev.com>

<https://linkedin.com/in/sergeyolontsev>



**Microsoft**  
**CERTIFIED**

Master

SQL Server® 2008

**Microsoft**  
**CERTIFIED**

Solutions Master

Charter - Data Platform

# Very Large Database (VLDB) Challenges

## *Why to refactor?*

Amount of Data


Transaction Rate

High Concurrency

Complex Logic

Number of Users

Developers / Hiring



# Before You Start






# Development Tools

You need a modern IDE + version control + CI/CD (Visual Studio + SSDT, IntelliJ DataGrip).

- Rename database objects / column names support
- Schema comparison / generate change script
- Better have dynamic SQL support / search in comments



# Tests



# SQL Server Features Which Helps Refactoring



1. Add Column (NULL, default value)
2. Views (partitioned views, updatable views)
3. Synonyms (point to another DB objects for decoupling, renaming objects)
4. Triggers (data sync, triggers on views)
5. Stored Procedures (mapping to another SP)
6. DML / DDL Transactions



# #1 Know Your Customer



# User Access Documentation

Do you have user access documentation in Excel or Confluence (Wiki)?



# User Access Documentation

Do you have user access documentation in Excel or Confluence (Wiki)?



Gathering information just before significant refactoring usually leads to fail!

# Document Interfaces Through Tests

1. Stored Procedures
2. Table Variables
3. Views
4. Direct Table Access



If someone has access to a table or user defined table variable, make a test, which will fail if a new column will be added, changed or removed.


# No Direct Access

Encapsulate data access. Stored procedures are a common way to encapsulate access to your database.

Decouple application code from database tables. Stored procedures are an effective way to decouple applications from database tables. They enable you to change database tables without changing application code.

Implement entity-based security access control (SAC). Instead of directly accessing source tables, applications instead invoke the relevant stored procedures.





## #2 Step-By-Step

# Step-By-Step Examples

## Merge Tables

1. Move column 1
2. Move column 2
3. ...
4. Move column N
5. Drop old table

# Step-By-Step Examples

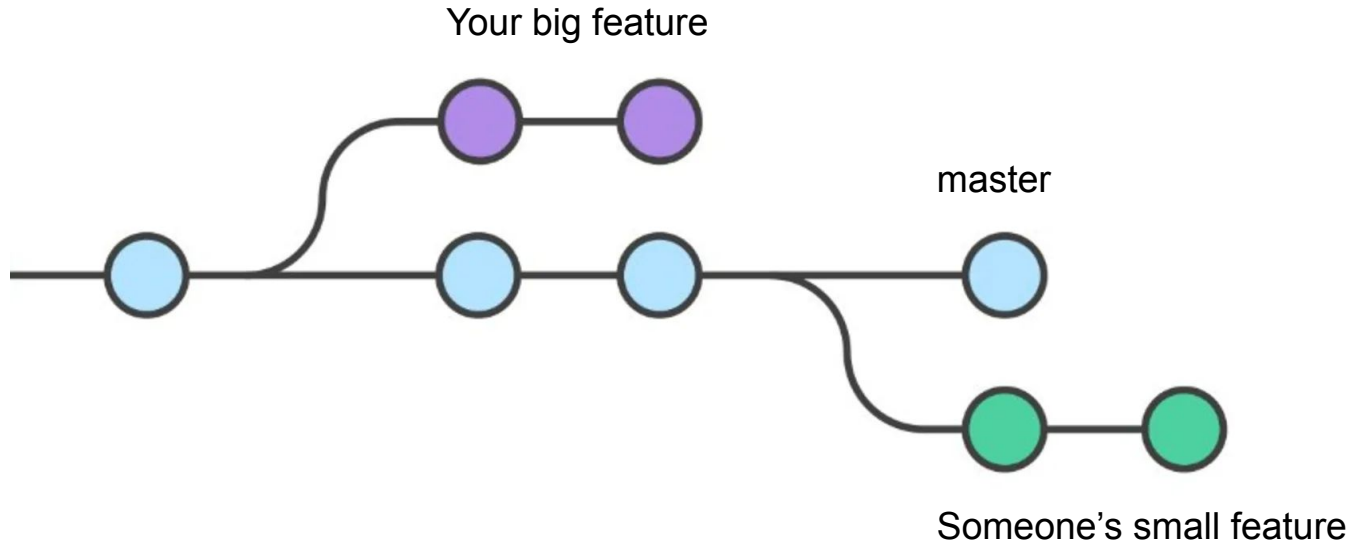
## Merge Tables

1. Move column 1
2. Move column 2
3. ...
4. Move column N
5. Drop old table

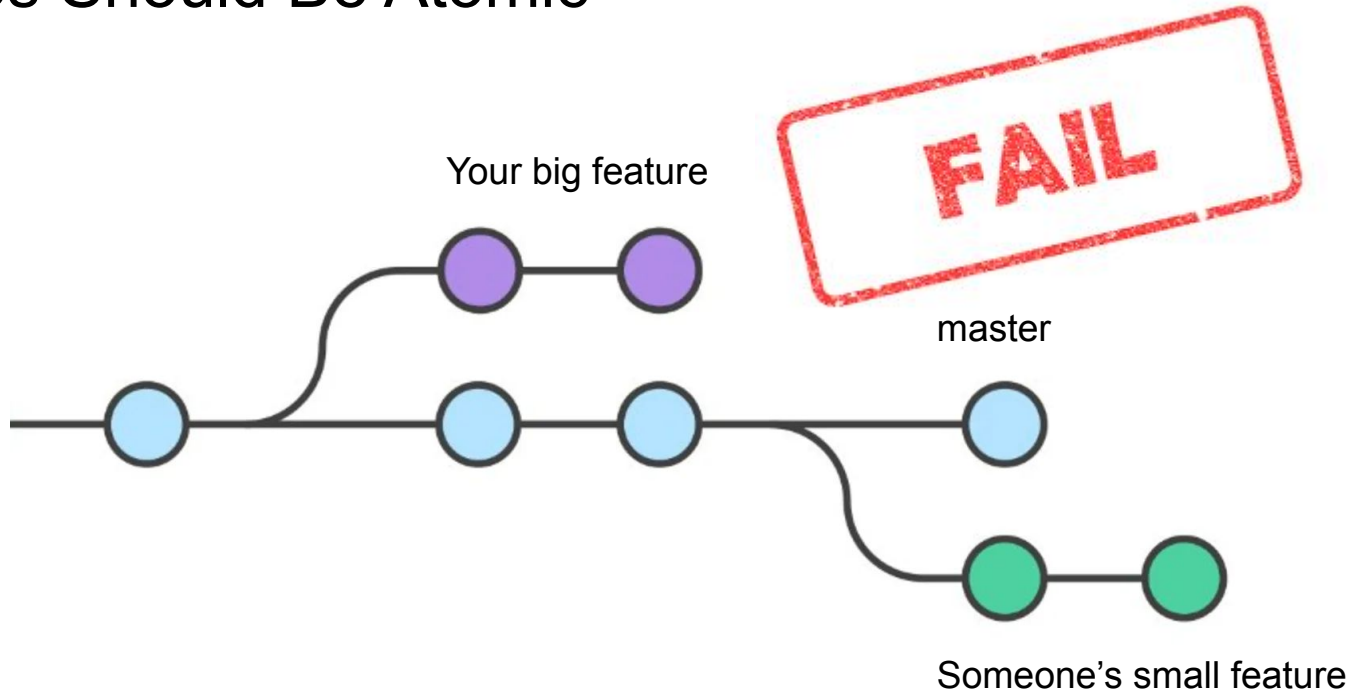
## Move Column


1. Create a new column
2. Move / sync data
3. Change queries
4. Drop old column

# Changes Should Be Atomic



# Changes Should Be Atomic

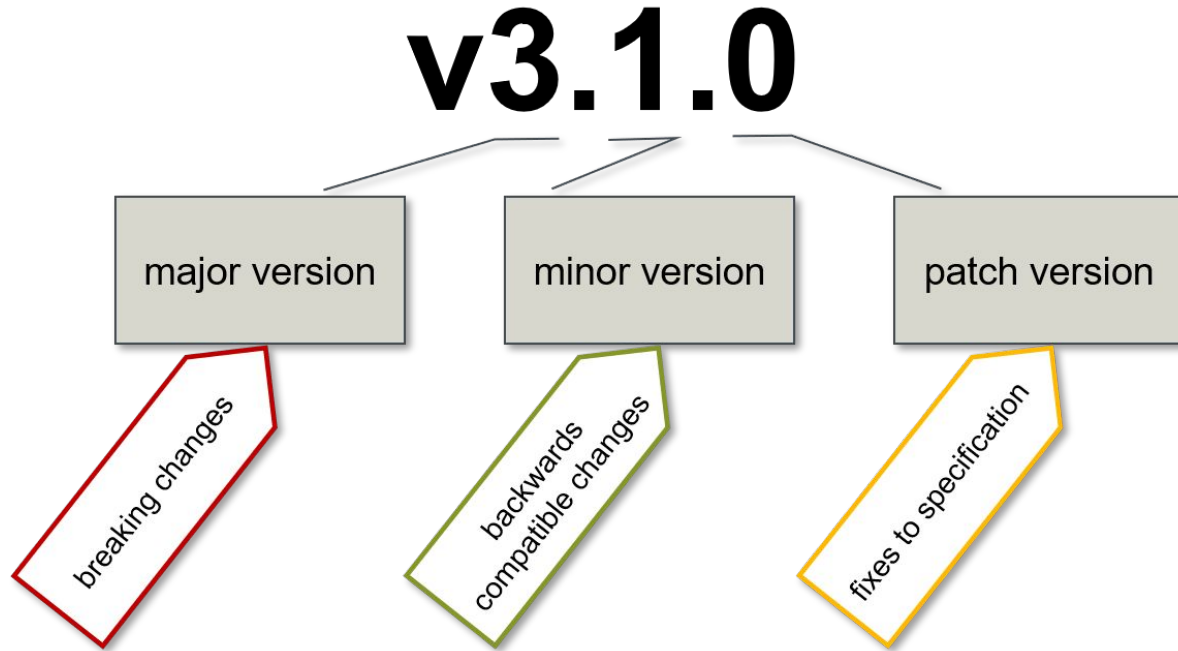




## #3 Deploy

Image Source: <https://octopus.com/blog/database-deployment-automation-adhoc-scripts>

# Deploy Scripts Versioning

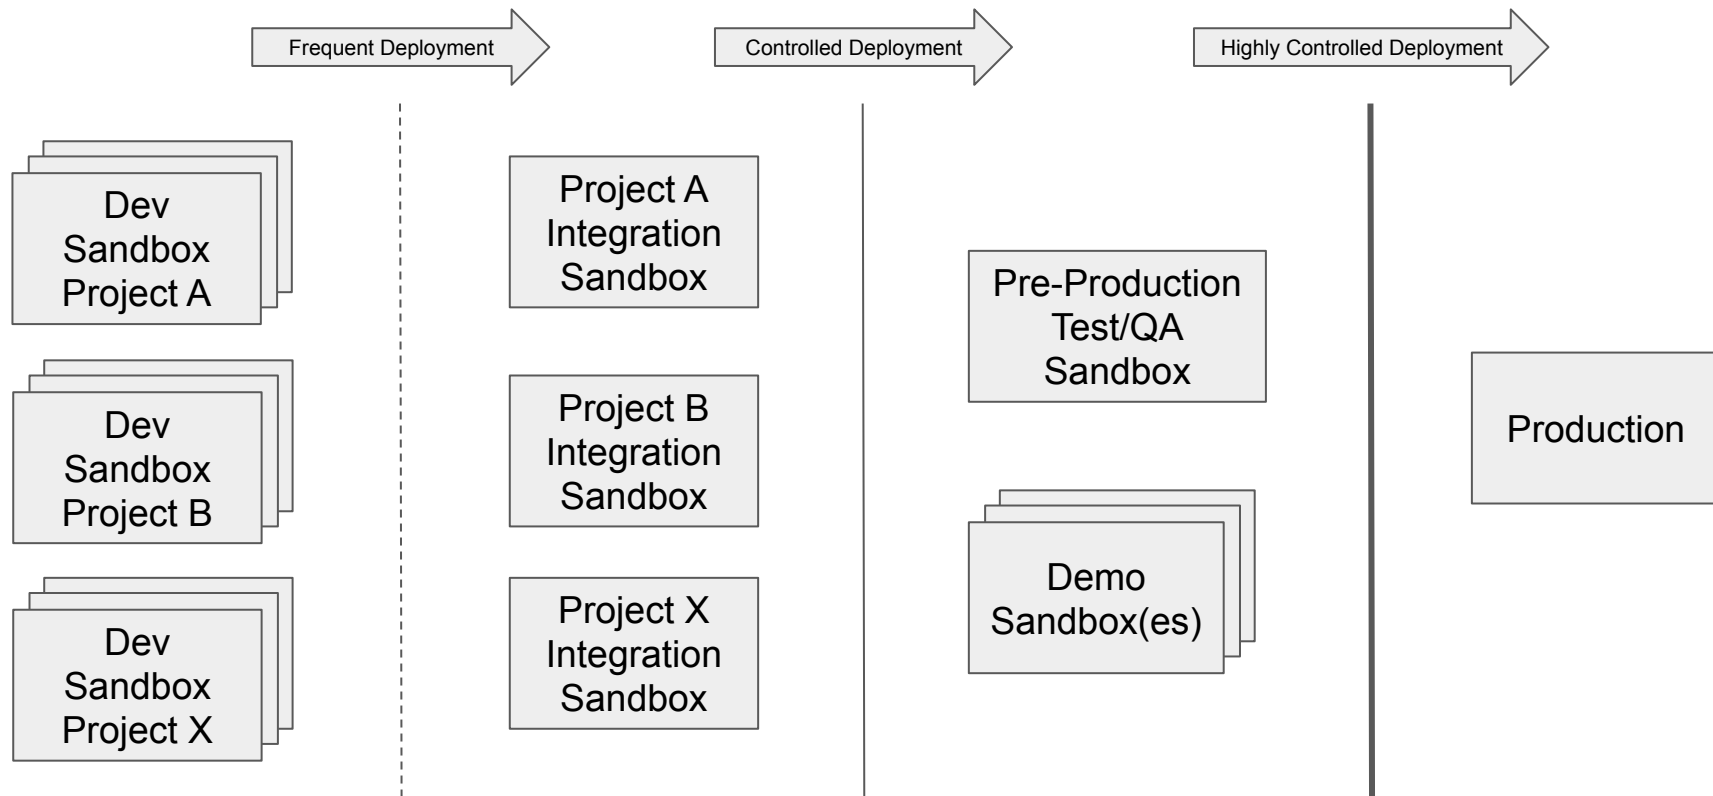


# Benefits of Deploy Scripts

1. Scripts are immutable.
2. If you have too many scripts, you can collapse old ones into one.
3. No need to deploy every script on production.
4. Easy to deploy to as many servers as needed.
5. Any application can easily add a DB structure inside it's projects of specific version. [A story comes...](#)



# Sandboxes




# Sandboxes

SQL Server in Docker was truly a game changer in software development!

It's fully customizable.

[https://hub.docker.com/\\_/microsoft-mssql-server](https://hub.docker.com/_/microsoft-mssql-server)






## #4 Backup and Recovery


# Backups / Restore / HA

## Piecemeal Restores


<https://learn.microsoft.com/en-us/sql/relational-databases/backup-restore/peicemeal-restores-sql-server?view=sql-server-ver16>



Full Restore




Step 1



Step 2

# A Queue Before Database



Does It Look Good?

delete from b

from [dbo].[SampleTable] as b

where b.[insert\_dt] < '2022-01-01'

# Long-Running Transactions

```
delete from b  
from [dbo].[SampleTable] as b  
where b.[insert_dt] < '2022-01-01'
```



# Controlled Batch Operations

```
declare @batch_size int = 10000  
while 1 = 1  
begin  
    delete top (@batch_size) from b  
    from [dbo].[SampleTable] as b  
    where b.[insert_dt] < '2022-01-01'  
  
    if @@rowcount = 0 break  
end
```







## #5 Move Column

Image Source: <https://allegiancemovingandstorage.com/>

# Move Unused Columns To Separate Table

| id | name | description      |
|----|------|------------------|
| 1  | Alex | Some Long Text 1 |
| 2  | Bob  | Some Long Text 2 |
| 3  | Mary | Some Long Text 3 |
| 4  | Kate | Some Long Text 4 |

Imagine: 90% of your queries do not use description column.

What about memory utilization?

# Move Unused Columns To Separate Table

| id | name |
|----|------|
| 1  | Alex |
| 2  | Bob  |
| 3  | Mary |
| 4  | Kate |

| id | description      |
|----|------------------|
| 1  | Some Long Text 1 |
| 2  | Some Long Text 2 |
| 3  | Some Long Text 3 |
| 4  | Some Long Text 4 |

## Memory Utilization Improvement


**10% → 11.11%**      **20% → 25%**      **30% → 42.86%**      **35% → 53.85%**

## Add Column → Join Elimination

| id | name | has_children |
|----|------|--------------|
| 1  | Alex | true         |
| 2  | Bob  | false        |
| 3  | Mary | false        |
| 4  | Kate | true         |

~2x performance increase.

Less cache is used for the second table.



## #6 Fragmentation

# Internal Fragmentation

## FILLFACTOR

<https://learn.microsoft.com/en-us/sql/relational-databases/indexes/specify-fill-factor-for-an-index?view=sql-server-ver16>

Page fill **70%** → **90%**, memory utilization increase **28.57%**


Page fill **60%** → **90%**, memory utilization increase **50%**

Page fill **55%** → **80%**, memory utilization increase **45.45%**

# External Fragmentation

There is one simple solution here...





## #7 Data Synchronization



# Data Migration / Synchronization Techniques

1. Triggers (on tables, on temp views)
2. Bulk Load
3. Bulk Load + Replay Recent Events
4. Parallel Run
5. Unified Data Load & Processing
  - a. Low Priority Queue for Bulk Load
  - b. Ability to overwrite existing data
6. Problems With Delete (ghost records, use PAGLOCK hint)

# DELETE WITH PAGLOCK

```
declare @batch_size int = 10000

while 1 = 1

begin

    delete top (@batch_size) from b with (paglock)

    from [dbo].[SampleTable] as b

    where b.[insert_dt] < '2022-01-01'

    if @@rowcount = 0 break

end
```


DOES DATA COLLECTION WASTE TIME AND RESOURCES?

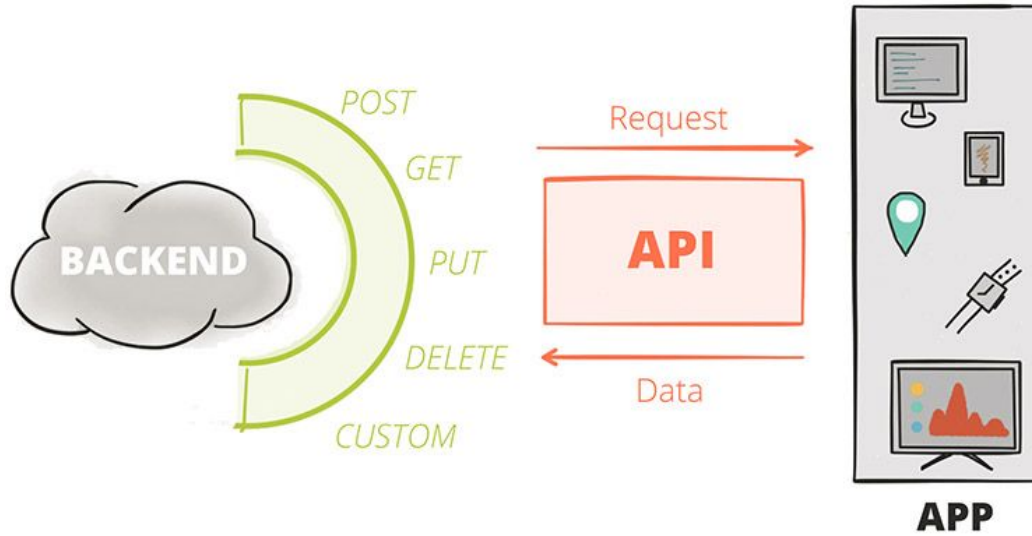


## #8 Stale Data

# Remove Stale Data & Structures

1. Unused Data
2. Unused Indexes
3. Unused Columns
4. Unused Tables





## #9 Application API

# External API Pros

1. Lookup tables in application (periodically refresh, no extra joins).
2. Benefit from external cache.
3. Not necessarily and application API.
  - a. Example: separate DB + Service Broker.
4. Easy to change DB engine.
5. Turn your database into a single application database.
  - a. Easier internal refactoring.
  - b. Split database.
  - c. Easy to change DB engine.





## #10 Data Integrity

# Constraints

Foreign key constraints reduce performance within your database because the existence of the row in the foreign table will be verified whenever the source row is updated.

If you drop foreign keys, you'll need better testing of your code and (probably) some regular external data checks.

Check constraints are good.







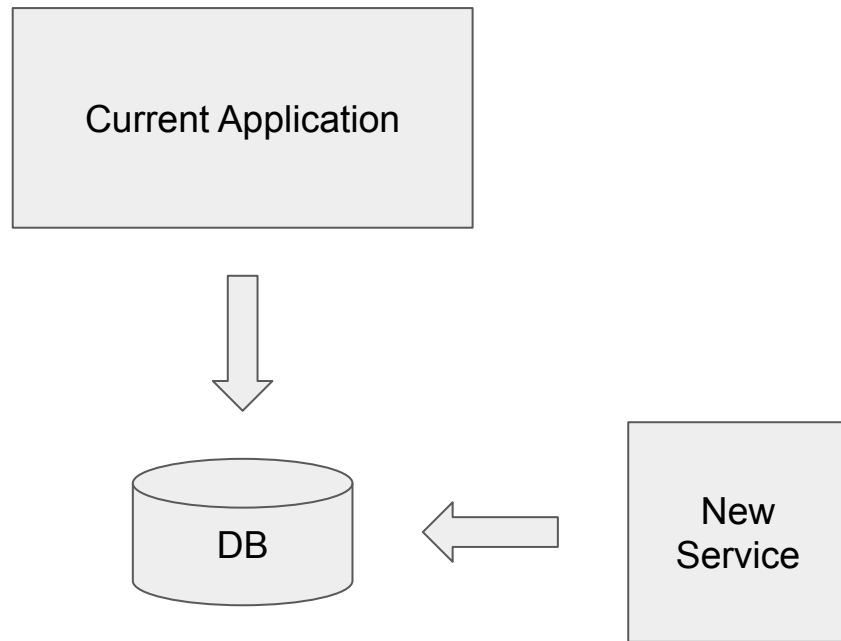

## #11 Split Database

Image Source: <https://diygarden.co.uk/log-splitters/best-log-splitting-axe>

# Split Database First











# Split Code First



# Split Database and Code Together

**NOT RECOMMENDED**

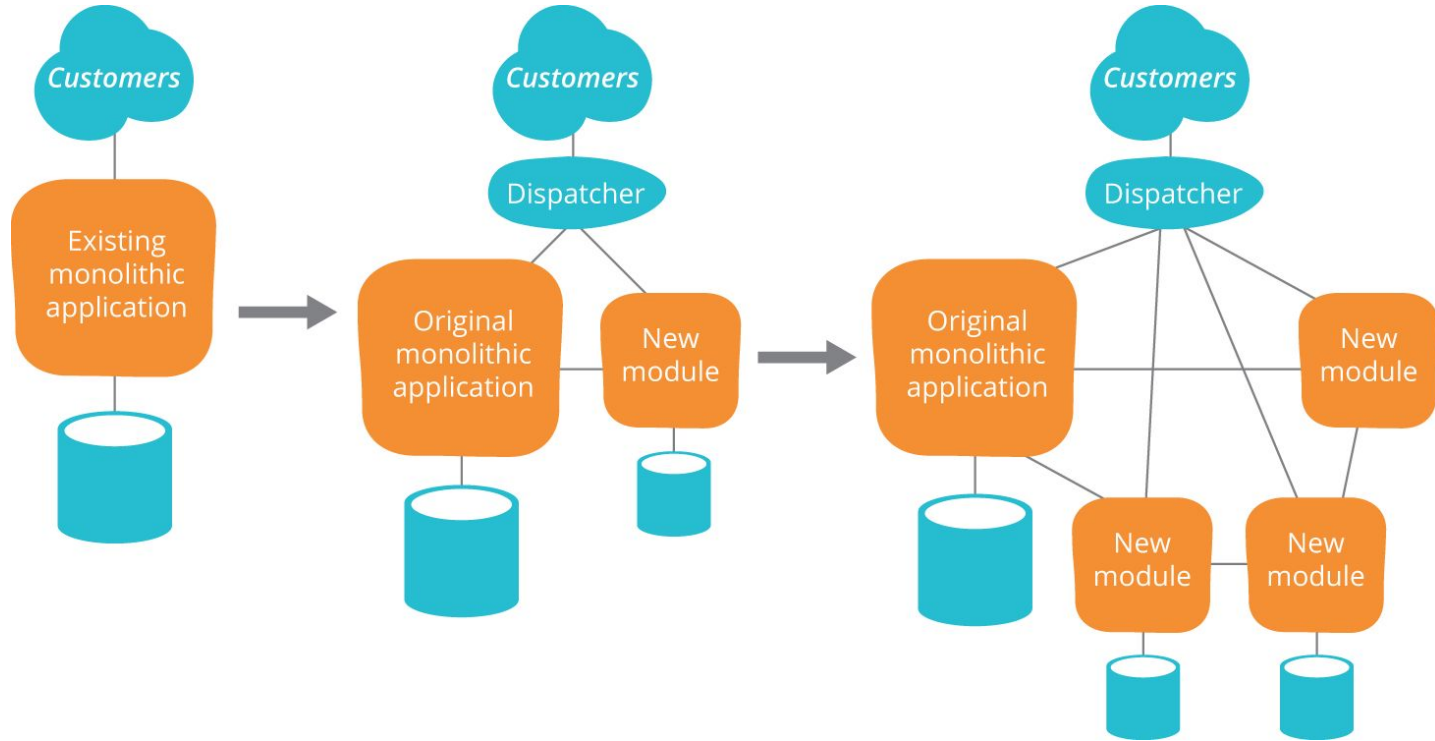


# Strangler Fig

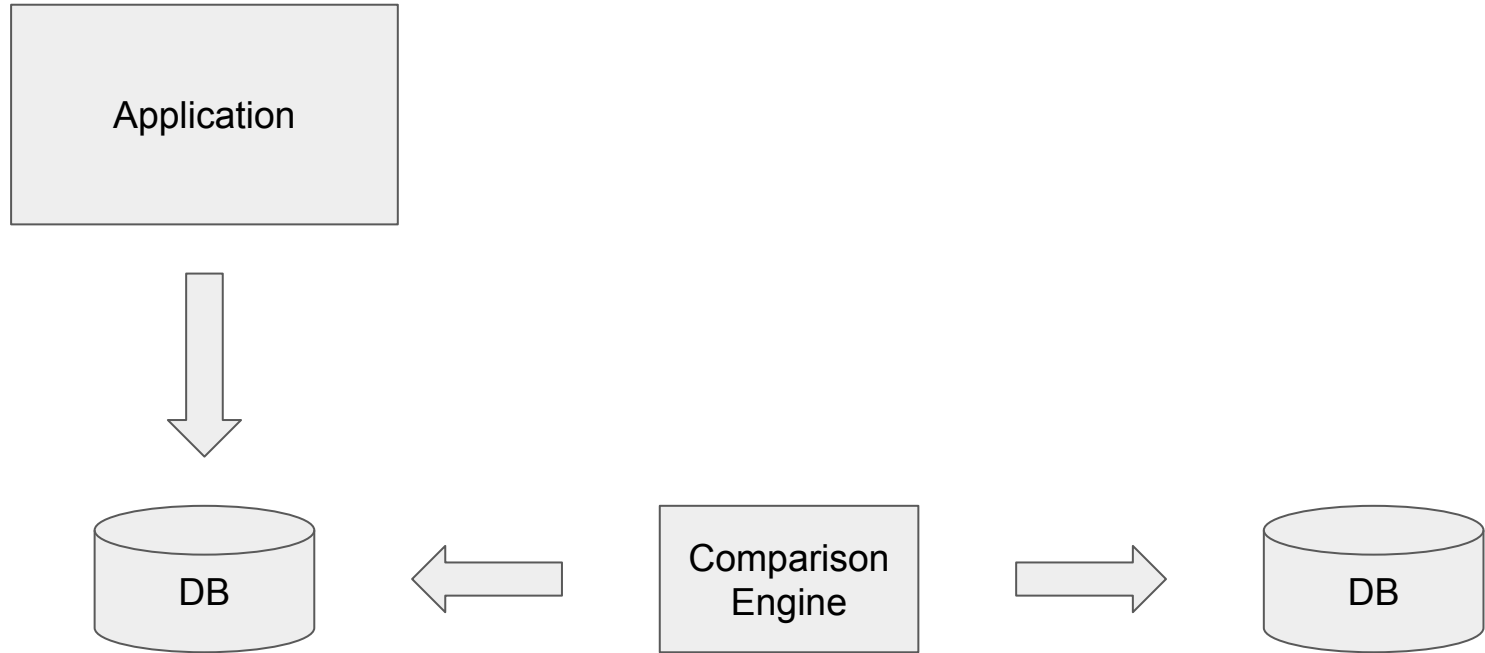
Incrementally migrate a legacy system by gradually replacing specific pieces of functionality with new applications and services. As features from the legacy system are replaced, the new system eventually replaces all of the old system's features, strangling the old system and allowing you to decommission it.

<https://learn.microsoft.com/en-us/azure/architecture/patterns/strangler-fig>

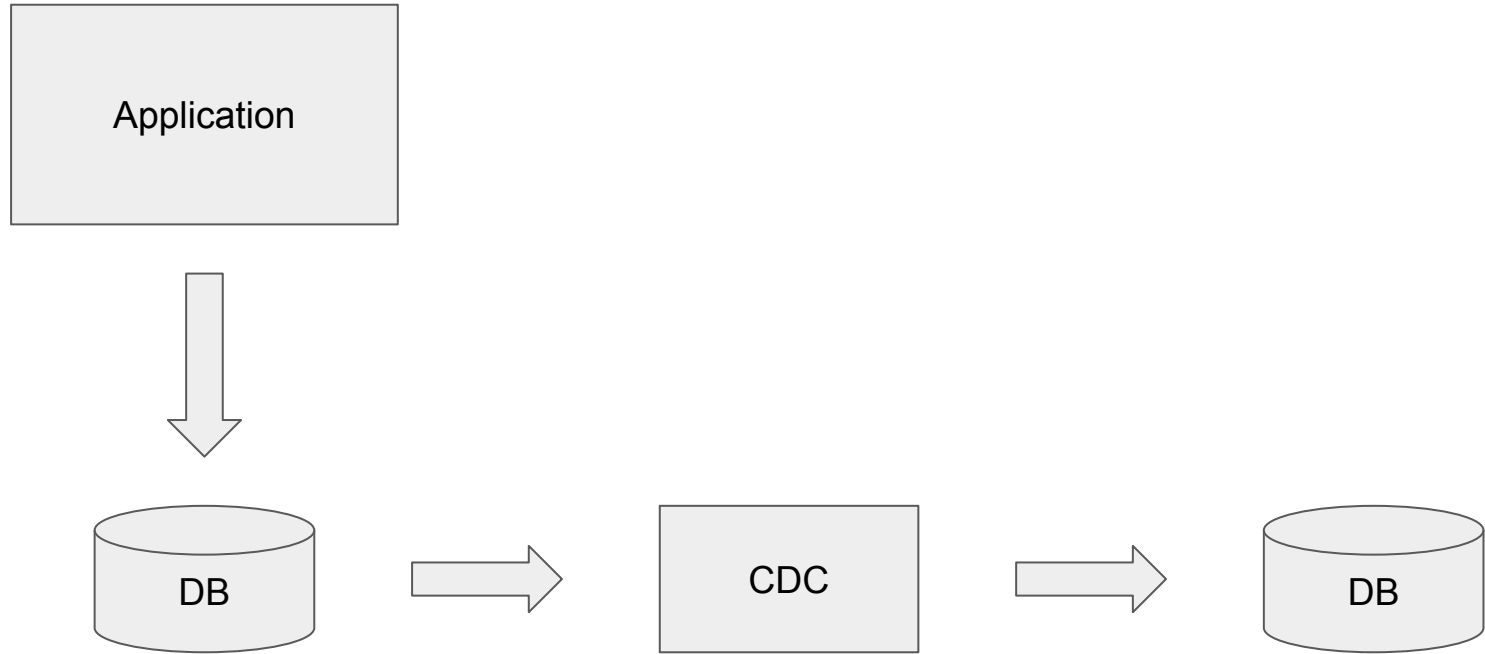
# Strangler Fig



# Parallel Run




# Change Data Capture





# Event Queue





# #1 Story: Optimization

# Replace One-To-Many With Associative Table

| user_id | name       |
|---------|------------|
| 1       | Alessandro |
| 2       | Luigi      |
| 3       | Isabella   |

| user_id | pizza_id |
|---------|----------|
| 1       | 1        |
| 1       | 2        |
| 2       | 2        |
| 3       | 1        |
| 3       | 2        |
| 3       | 3        |

| pizza_id | name                |
|----------|---------------------|
| 1        | Margherita          |
| 2        | Pepperoni           |
| 3        | Quattro<br>Formaggi |
| 4        | Parma               |

# Replace One-To-Many With Associative Table

| user_id | name       |
|---------|------------|
| 1       | Alessandro |
| 2       | Luigi      |
| 3       | Isabella   |

| set_id | pizzas    |
|--------|-----------|
| 1      | [1, 2]    |
| 2      | [2]       |
| 3      | [1, 2, 3] |

| pizza_id | name             |
|----------|------------------|
| 1        | Margherita       |
| 2        | Pepperoni        |
| 3        | Quattro Formaggi |
| 4        | Parma            |

| user_id | pizza_set_id |
|---------|--------------|
| 1       | 1            |
| 2       | 2            |
| 3       | 3            |



hard


**VS**



soft

## #2 Story: Hard vs. Soft Deletes

# Given...





## #3 Story: Plan In Advance

Image Source: [https://valley.d51schools.org/counseling/life\\_after\\_valley/plan\\_ahead](https://valley.d51schools.org/counseling/life_after_valley/plan_ahead)

# Given...

DB ~3 TB

Need to predict growth and plan resources.

3 tables each at about 600-800 GB. Fragmentation ~60-70%.

Fillfactor 100%. Free space in DB ~400GB.

Index reorganize? Index Rebuild? Rebuild by partition?




# Additional Resources

## *Refactoring Databases: Evolutionary Database Design*

Scott W. Ambler (Author), Pramod J. Sadalage (Author), Martin Fowler (Foreword), John Graham (Foreword), Sachin Rekhi (Foreword), Paul Dorsey (Foreword)

<https://www.amazon.com/gp/product/0321293533/>



The logo for Data Saturdays. It features the word "DATA" in a large, bold, dark blue sans-serif font. To the left of the "D" is a blue icon consisting of three horizontal bars of increasing length, resembling a stylized "D" or a data bar chart. Below "DATA" is the word "SATURDAYS" in a smaller, blue, spaced-out sans-serif font.

DATA  
SATURDAYS

*Crazie*

The logo for engage LABS. The word "engage" is written in a lowercase, rounded sans-serif font. The "e" at the start is dark blue and contains a white circular graphic with radiating lines. The "n" is dark blue, and the "g" is a lighter blue. The "a" is a light green, and the "c" is a slightly darker green. The "e" at the end is a light green and also contains a white circular graphic with radiating lines. Below "engage" is the word "LABS" in a small, dark blue, spaced-out sans-serif font.

engage  
LABS